

Double-click (or enter) to edit

```
import numpy as np
```

```
import pandas as pd
```

```
import matplotlib.pyplot as plt
import seaborn as sns
```

```
df = pd.read_csv(r"https://github.com/YBI-Foundation/Dataset/raw/main/Boston.csv")
```

```
df.head()
```

	CRIM	ZN	INDUS	CHAS	NX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0	0.538	6.575	65.2	4.0900	1	296.0	15.3	396.90	
1	0.02731	0.0	7.07	0	0.469	6.421	78.9	4.9671	2	242.0	17.8	396.90	
2	0.02729	0.0	7.07	0	0.469	7.185	61.1	4.9671	2	242.0	17.8	392.83	
3	0.03237	0.0	2.18	0	0.458	6.998	45.8	6.0622	3	222.0	18.7	394.63	
4	0.06905	0.0	2.18	0	0.458	7.147	54.2	6.0622	3	222.0	18.7	396.90	

```
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 506 entries, 0 to 505
Data columns (total 14 columns):
#   Column      Non-Null Count  Dtype
---  -
0    CRIM        506 non-null    float64
1    ZN          506 non-null    float64
2    INDUS       506 non-null    float64
3    CHAS        506 non-null    int64
4    NX          506 non-null    float64
5    RM          506 non-null    float64
6    AGE         506 non-null    float64
7    DIS         506 non-null    float64
8    RAD         506 non-null    int64
9    TAX         506 non-null    float64
10   PTRATIO     506 non-null    float64
11   B           506 non-null    float64
12   LSTAT       506 non-null    float64
13   MEDV        506 non-null    float64
dtypes: float64(12), int64(2)
memory usage: 55.5 KB
```

```
df.describe()
```

	CRIM	ZN	INDUS	CHAS	NX	RM	AGE
count	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000	506.000000
mean	3.613524	11.363636	11.136779	0.069170	0.554695	6.284634	68.574900
std	8.601545	23.322453	6.860353	0.253994	0.115878	0.702617	28.148800
min	0.006320	0.000000	0.460000	0.000000	0.385000	3.561000	2.900000
25%	0.082045	0.000000	5.190000	0.000000	0.449000	5.885500	45.025000
50%	0.256510	0.000000	9.690000	0.000000	0.538000	6.208500	77.500000
75%	3.677083	12.500000	18.100000	0.000000	0.624000	6.623500	94.075000
max	88.976200	100.000000	27.740000	1.000000	0.871000	8.780000	100.000000



```
df.columns
```

```
Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',
      'PTRATIO', 'B', 'LSTAT', 'MEDV'],
      dtype='object')
```

```
y = df["MEDV"]
```

```
X = df[['CRIM', 'ZN', 'INDUS', 'CHAS', 'NX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',
      'PTRATIO', 'B', 'LSTAT']]
```

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test = train_test_split(X,y,test_size = 0.3,random_state =
```

```
X_train.shape,X_test.shape,y_train.shape,y_test.shape
```

```
((354, 13), (152, 13), (354,), (152,))
```

```
from sklearn.preprocessing import StandardScaler
```

```
sc = StandardScaler()
```

```
X_train_sc = sc.fit_transform(X_train,y)
```

```
X_test_sc = sc.fit_transform(X_test,y)
```

```
X_train_sc
```

```
array([[ -0.14113619, -0.48175769, -0.19860022, ...,  0.00438903,
        -0.05084503, -0.01555641],
       [ -0.42121529,  3.02166196, -1.33410259, ..., -1.68641979,
        0.42969249, -1.33650784],
       [ -0.41266839, -0.48175769,  0.22414717, ...,  0.14148164,
        0.19739169, -0.10842497],
       ...,
       [ -0.38944304, -0.48175769, -0.19860022, ...,  0.00438903,
        0.37963873,  0.77313338],
       [ -0.41404001,  0.41002186, -0.81324318, ..., -0.72677154,
        0.43161763,  0.09671754],
       [ -0.41578561,  2.06618387, -1.3831586 , ..., -0.04130851,
        0.39707198, -0.68781395]])
```

X_test_sc

```
array([[ -0.36714008, -0.50235603, -0.6925381 , ..., -0.57641511,
        0.2366856 , -1.24860568],
       [ -0.40880876, -0.50235603, -0.58591169, ..., -0.33768188,
        0.43031542, -0.31886558],
       [ -0.41291768, -0.50235603, -0.12035979, ..., -0.38542852,
        0.36717526,  0.17122998],
       ...,
       [ -0.34428827, -0.50235603,  1.66375525, ...,  1.23795746,
        0.30005961, -0.18769294],
       [ -0.05769974, -0.50235603,  1.31684399, ..., -1.86557456,
        -0.3514533 , -0.15886379],
       [ -0.42293258,  1.25907688, -0.66100071, ..., -0.48092181,
        0.43031542, -0.75418575]])
```

```
from sklearn.neighbors import KNeighborsRegressor
```

```
knnreg = KNeighborsRegressor(n_neighbors=3)
```

```
knnreg.fit(X_train_sc,y_train)
```

```
KNeighborsRegressor(n_neighbors=3)
```

```
y_pred = knnreg.predict(X_test_sc)
```

```
from sklearn.metrics import mean_absolute_error,mean_absolute_percentage_error
```

```
mean_absolute_error(y_test,y_pred)
```

```
2.5309210526315793
```

```
mean_absolute_percentage_error(y_test,y_pred)
```

```
0.12289250095895413
```


```
error_rate = []

for i in range(1,40):
    knn = KNeighborsRegressor(n_neighbors=i)
    knn.fit(X_train_sc,y_train)
    pred_i = knn.predict(X_test_sc)
    error_rate.append(mean_absolute_error(y_test,pred_i))
```

error_rate

```
[3.134210526315789,
 2.795065789473684,
 2.5309210526315793,
 2.5592105263157894,
 2.672894736842105,
 2.6517543859649124,
 2.7910714285714286,
 2.8222039473684206,
 2.8627192982456138,
 2.8897368421052634,
 2.9014952153110047,
 2.9251644736842106,
 2.9426619433198375,
 2.87687969924812,
 2.9153070175438596,
 2.943215460526315,
 2.928405572755418,
 2.9784722222222215,
 2.960560941828255,
 2.972763157894736,
 2.9944235588972425,
 3.029186602870813,
 3.02883295194508,
 3.0328673245614035,
 3.0655,
 3.0812500000000003,
 3.080750487329435,
 3.117763157894737,
 3.128493647912886,
 3.137214912280702,
 3.133000848896435,
 3.1576685855263156,
 3.1555223285486447,
 3.1607972136222906,
 3.1681766917293235,
 3.1903508771929823,
 3.192247510668563,
 3.2162049861495845,
 3.2415485829959514]
```

```
fig,ax= plt.subplots(figsize = (12,5))
ax.plot(range(1,40),error_rate,color = 'Red')
plt.show()
```



✓ 0s completed at 9:00 AM

